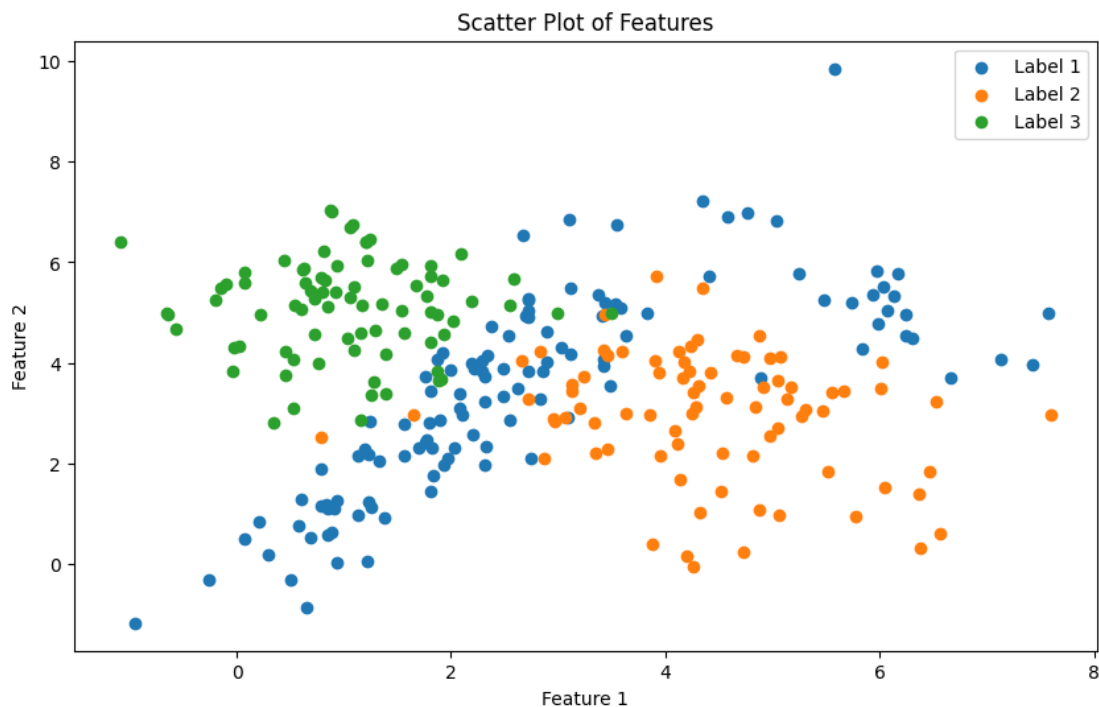


# Pattern Recognition

Γιακουμόγλου Πασχάλης 10054

Γκύζης Γεώργιος 10107

# Data Visualization and Split (Part A-C)



- ▶ Clear distribution of data points across two features.
- ▶ Overlapping regions suggest potential challenges for classification.
- ▶ Class 1: 42.86%, Class 2: 28.57%, Class 3: 28.57% (No heavy imbalances)
- ▶ Similar values for both features (no need for scaling,  $\text{var}(\text{feat1}) \approx \text{var}(\text{feat2})$ )
- ▶ The data will be split 0.5-0.5
- ▶ The same datasets will be used for Parts A,B,C

# Classifier Regions (Part A-C)

- ▶ Mesh Grid that covers the latent space of the features
- ▶ Classify each point of the grid
- ▶ Different coloring for each region and different shapes for training and test data

# Part A: Bayes Gaussian Classifier

- ▶ The classifier assumes gaussian distribution of the features of each class
- ▶ Two cases are implemented: Same or different covariance matrices for each class
- ▶ Same covariance matrix: simpler model, danger of underfitting
- ▶ Different covariance matrix: complex model, danger of overfitting but can capture more complex patterns

# Part A: Bayes Gaussian Classifier

- ▶ For the distribution estimation the maximum likelihood technique is followed
- ▶ For the Gaussian Distribution, the maximum likelihood estimators are:

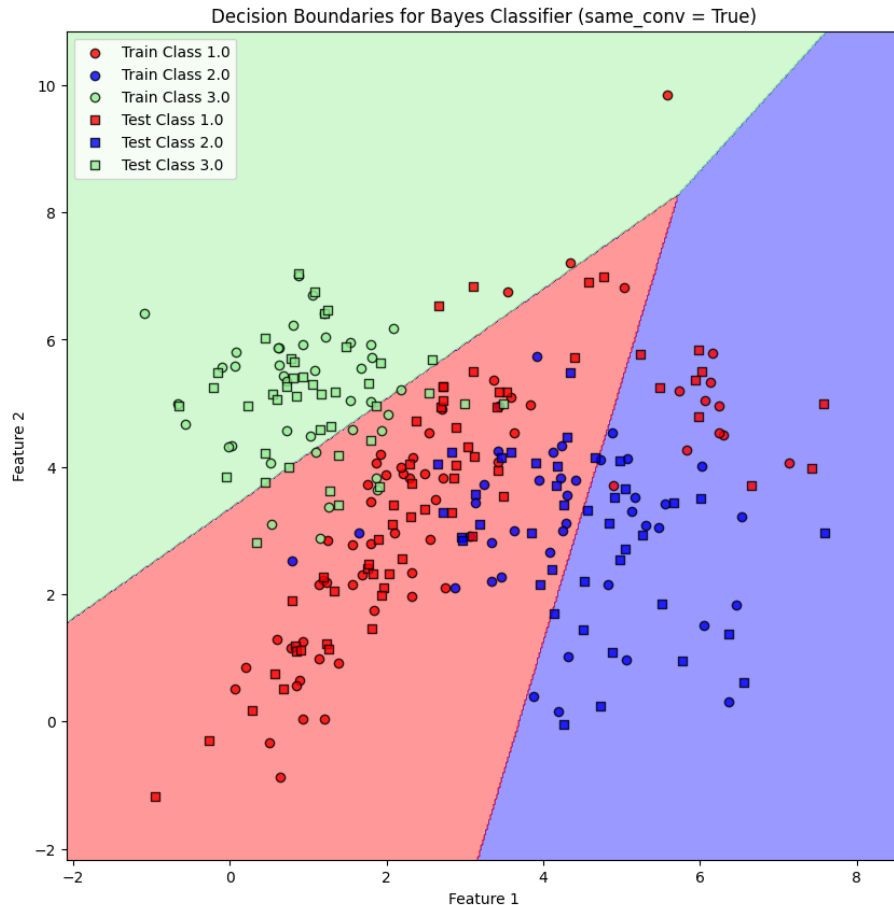
$$\hat{\mu}_n = \frac{1}{n} \sum_{j=1}^n x_j$$

$$\hat{V}_n = \frac{1}{n} \sum_{j=1}^n (x_j - \hat{\mu}_n)(x_j - \hat{\mu}_n)^T$$

# Part A: Bayes Gaussian Classifier

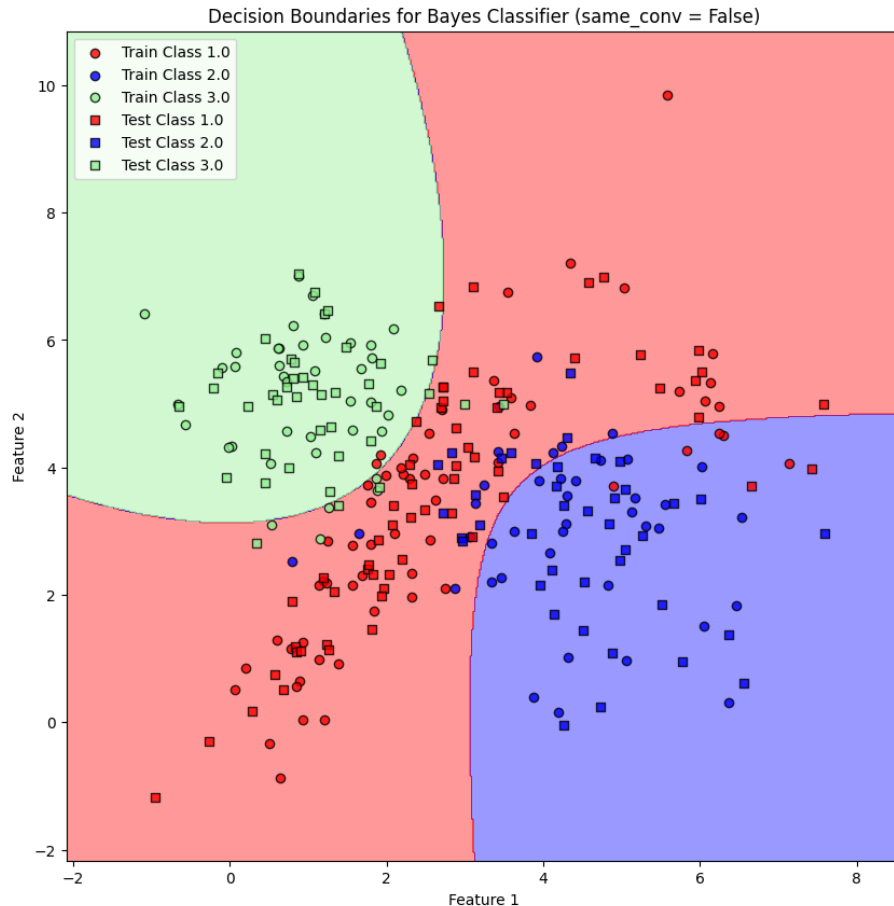
- ▶ In our implementation the classifier has similar methods to the sklearn (fit, predict, predict\_proba)
- ▶ That ensures that our estimator could be used later with other methods of sklearn (cross\_val, ensembles, etc.)

# Part A: Bayes Gaussian Classifier



- ▶ Assuming the same covariance matrix, the classifier achieves 71% accuracy on the test set
- ▶ The same covariance matrix creates linear boundaries for each region
- ▶ The classifier is limited by the linear boundaries
- ▶ Cluster of points of the same class are incorrectly classified
- ▶ More complex classifier could potentially capture these patterns

# Part A: Bayes Gaussian Classifier



- ▶ Assuming a different covariance matrix, the classifier achieves 85% accuracy on the test data
- ▶ The different covariance matrix creates non linear boundaries for each region
- ▶ The classifier can capture more intricate patterns in the data
- ▶ This is confirmed by the improved accuracy on the training set



# Part B: K-Nearest Neighbors Classifier

- ▶ sk-learn implementation of the classifier
- ▶ Euclidean distance as a distance metric
- ▶ Lower of values of k-NN are prone to overfitting
- ▶ Higher values of k-NN usually can not capture complex patterns in the dataset

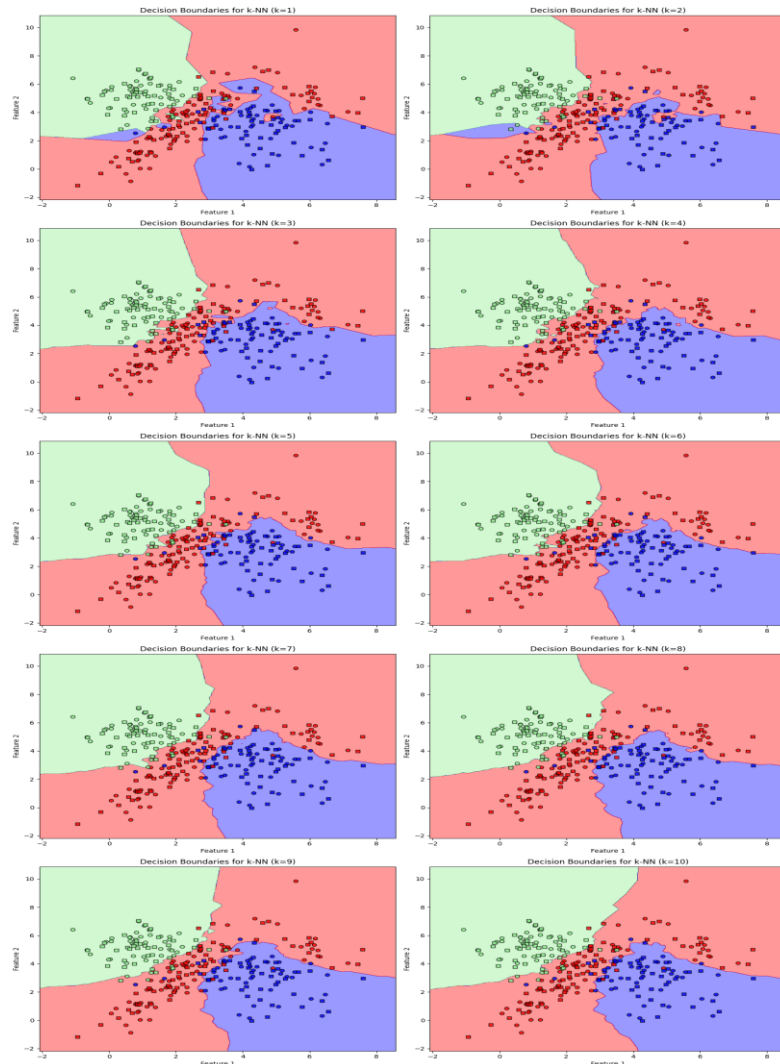
# Part B: K-Nearest Neighbors Classifier

## Results of the K-NN

Number of neighbor	Accuracy (Test Set)
1-NN	88.57%
2-NN	87.14%
3-NN	87.14%
4-NN	87.14%
5-NN	89.28%
6-NN	87.58%
7-NN	88.58%
8-NN	89.28%
9-NN	90.00%
10-NN	88.00%

- ▶ K=1: Complex decision boundary, likely overfitting.
- ▶ Slight decrease then plateau around 87.14% for K = 2 to K = 4.
- ▶ Increase to 89.23% at K = 5
- ▶ High but fluctuating trend, peaking at 90% for k=9

# Part B: K-Nearest Neighbors Classifier



- ▶  $K=1$ : Complex decision boundary, likely overfitting.
- ▶  $K = 3$  to  $K = 5$ : Smoother boundaries indicate improved generalization.
- ▶  $K = 6$  to  $K = 10$ : Continued smoothness with effective class separation, balanced bias and variance.
- ▶ Higher  $K$  values show no underfitting, maintaining model effectiveness.

# Bayes vs KNN

- ▶ Comparing the accuracy scores, the k-NN classifier's best result (90% for  $k=9$ ) outperforms the Gaussian Bayes classifier's accuracy of 85% with separate covariance matrices. This suggests that k-NN may be more adept at handling this particular dataset's intricacies without making strong distributional assumptions.

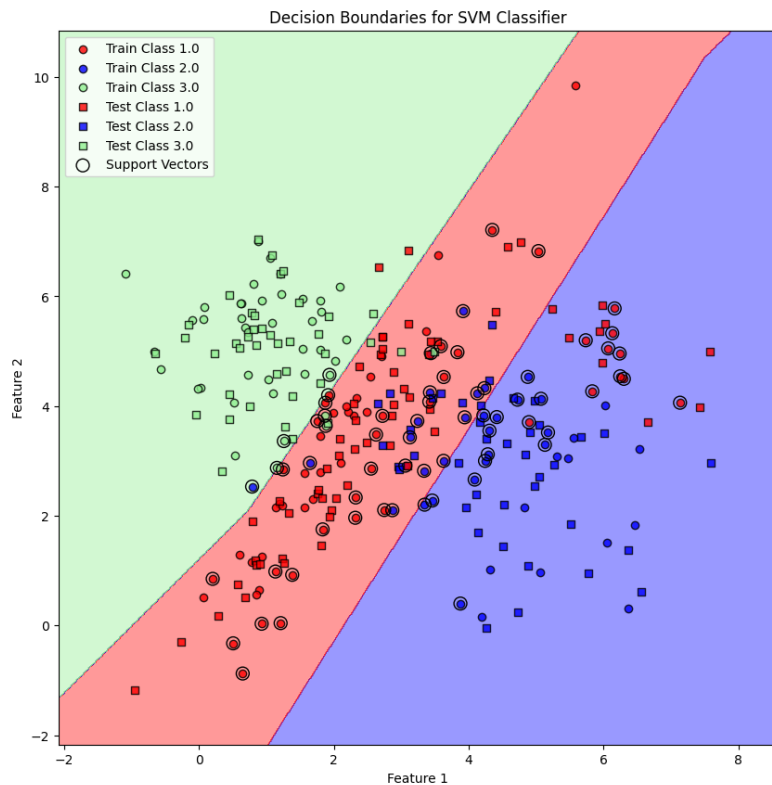
# Part C: Linear SVMs

- ▶ Train a linear SVM classifier with GridSearch for optimal 'C'.
- ▶ 'C' parameter balances classifier simplicity and training performance.
- ▶ Higher 'C' values increase model complexity and cost of misclassification.
- ▶ Aim to find 'C' that fits well without causing overfitting.

# Part C: Linear SVMs

- ▶ For the grid search we perform cross validation on the train set and use the model with the best average accuracy on the test set
- ▶ Best C: 100
- ▶ Accuracy on the test set: 0.8

# Part C: Linear SVMs



- ▶ Notable class overlap and dense support vectors for classes 2 and 3.
- ▶ Linear SVM outperforms Bayes model with same covariance matrix.
- ▶ Class overlap hints at linear SVM's limitation in capturing data complexity.
- ▶ Suggests potential benefits of using non-linear kernels for complexity modeling.

# Part C: SVM with RBF kernel

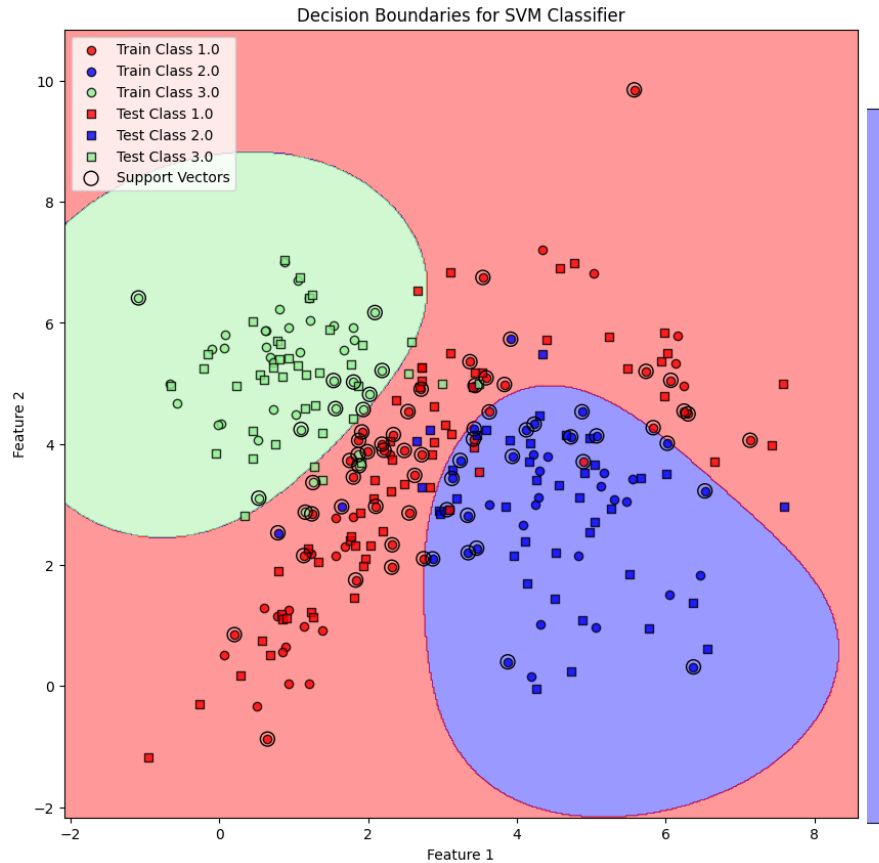
- ▶ Perform GridSearch on SVM with RBF kernel for 'C' and 'gamma' hyperparameters.
- ▶ 'Gamma' determines the influence range of support vectors.
- ▶ High 'gamma': small influence radius, irregular boundaries, risk of overfitting.
- ▶ Low 'gamma': large influence radius, smoother boundaries, risk of underfitting.



# Part C: SVM with RBF kernel

- ▶ For the grid search we perform cross validation on the train set and use the model with the best average accuracy on the test set
- ▶ Best score cv: 90%
- ▶ Best C: 1
- ▶ Best  $\gamma$ : 0.1
- ▶ Accuracy on the test set: 88.57%

# Part C: SVM with RBF kernel



- ▶ Highly Complex decision Regions
- ▶ Well-tuned model with balanced decision boundaries.
- ▶ rbf-SVM vastly outperforms linear model.
- ▶ Fewer support vectors than linear SVM, capturing complex patterns.

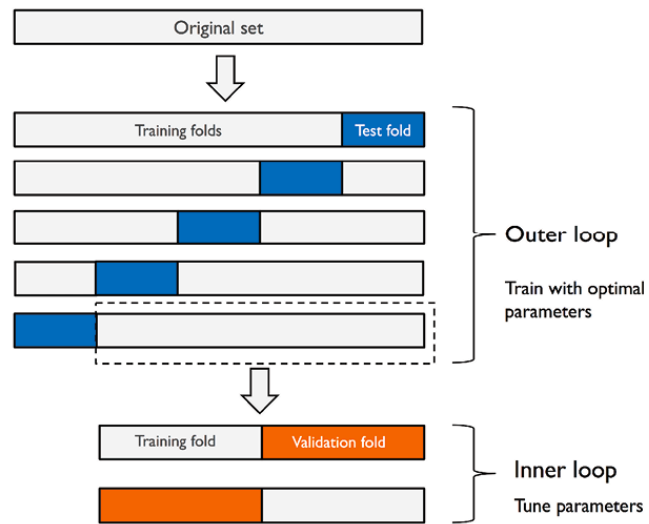
# Overall

- ▶ Best k-NN classifier leads with 90% test accuracy.
- ▶ RBF kernel SVM follows closely at 88.57%.
- ▶ Gaussian Bayes with separate covariance matrices: 85% accuracy.
- ▶ Linear SVM shows lower performance at 80%.
- ▶ Gaussian Bayes with shared covariance matrix has the least at 72.42%.

# Part D: Model Selection

- ▶ Cross-Validation Practice:
  - ▶ Perform cross-validation on training data.
  - ▶ Select model with best cross-validation accuracy.
  - ▶ Model that best fits in each validation fold, not model with least generalization error
- ▶ Nested Cross-Validation for Generalization:
  - ▶ Use nested cross-validation to estimate generalization error.
  - ▶ Outer loop divides data into training and testing folds.
  - ▶ Inner loop performs cross-validation on each training fold.
  - ▶ Average accuracy reported on outer testing sets.

# Part D: Model Selection



- ▶ Purpose of Nested Cross-Validation:
  - ▶ Utilized for model selection, not hyperparameter tuning.
  - ▶ After model selection, conduct cross-validation on full training set for optimal hyperparameters.

<https://vitalflux.com/python-nested-cross-validation-algorithm-selection/>

# Part D: Bayesian Optimization

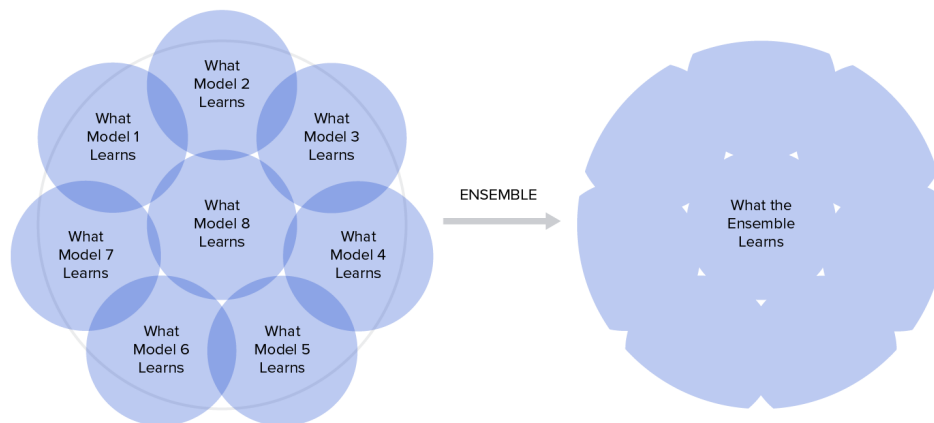
- ▶ Bayesian Optimization Overview:
  - ▶ Sequential strategy for globally optimizing expensive-to-evaluate black-box functions.
  - ▶ Treats objective functions as random function, using priors for initial behavior assumptions.
  - ▶ Updates priors to posteriors post-evaluation, guiding next query points.
- ▶ Application in Classifier Optimization:
  - ▶ Applied to classifiers as black-box functions.
  - ▶ Uses 5-fold cross-validation accuracy as the evaluation metric.
- ▶ Implementation Tool:
  - ▶ Utilized the scikit-optimize library for implementation.

# Part D: Bayesian Vs Grid Search

- ▶ **Efficiency:** Bayesian optimization is more efficient as it requires fewer function evaluations. It intelligently chooses the next point to evaluate based on past evaluations.
- ▶ **Handling Complex Spaces:** Better suited for optimizing over high-dimensional and complex spaces where grid search becomes impractical.
- ▶ **Flexible and Informative:** Adjusts exploration and exploitation dynamically, making it more informative for decision-making compared to the exhaustive nature of grid search.

# Part D: Ensembles

- ▶ Ensembles:
  - ▶ Combine multiple models to enhance prediction accuracy.
  - ▶ Leverage strengths and mitigate weaknesses of individual models.
  - ▶ Most Decorated Model Family in Kaggle competitions



<https://www.toptal.com/machine-learning/ensemble-methods-kaggle-machine-learn>



# Part D: Boosting

- ▶ Boosting Overview:
  - ▶ Sequentially trains models, focusing more on instances misclassified by previous models.
  - ▶ Weighs each model and instance to prioritize correcting errors, enhancing overall performance.
- ▶ AdaBoost
- ▶ XGBoost
  - ▶ Solves hyperparameter tuning complexity in boosting by optimizing a loss function across all models.
  - ▶ Known for efficiency and performance, incorporating techniques like parallel computing and regularization. Won recognition for its effectiveness in a Kaggle competition.

# Part D: Bagging

- ▶ Bagging Overview:
  - ▶ Utilizes bootstrapping to create overlapping data subsets.
  - ▶ Aggregates models' predictions using mean, median, mode, or weighted metrics.
- ▶ RandomForest:
  - ▶ Applies bagging with decision trees, using random subsets of attributes at each node.
  - ▶ Ensures diversity among trees, enhancing performance.
  - ▶ Effective for feature selection, offers good performance with low variance and training time.

# Part D: Voting Classifiers

- ▶ Voting Classifier Concept:
  - ▶ Combines predictions from multiple models.
  - ▶ Utilizes either 'hard' voting (majority rule) or 'soft' voting (weighted probabilities)
- ▶ Advantages:
  - ▶ Leverages strengths of various models.
  - ▶ Can improve prediction robustness and accuracy.

# Part D: Stacking Classifiers

- ▶ Stacking Overview:
  - ▶ Combines multiple models' predictions as input for a final meta-model.
  - ▶ Enhances predictive accuracy beyond individual models.
- ▶ Training Methodology:
  - ▶ Base models are trained on a part of the dataset.
  - ▶ Their out-of-sample predictions form a new dataset to train the meta-model.
  - ▶ This methodology ensures that both the training and testing dataset come from the same distribution or the the meta classifier doesn't prefer the most overfitted model

# Part D: Stacking Classifiers

- ▶ Limitations of the sk-learn implementation:
  - ▶ Can not use engineered features as input to the meta classifier
  - ▶ Computational inefficiency: If we want to evaluate base models we must retrain them
  - ▶ Difficult to fine tune: The more base models are used the more difficult it is to fine tune them

# Part D: Final Model

- ▶ To select a model we used an outer loop of 2 folds and an inner loop of 5 folds
- ▶ The following models were tried: SVM (Bayes search), MLP (Grid search), 1d-CNN (Grid search), XGBoost (Bayes search), RandomForest (Bayes search), HistGradBoosting (bayes search), Stacking (Grid Search), AdaBoost (Bayes search)
- ▶ Our winner was a stacking classifier with 10 base models and a soft voting classifier with (xgb, MLP, adaboost, SVM) as a meta classifier. It achieved an accuracy of 86% in the outer folds.

# Results

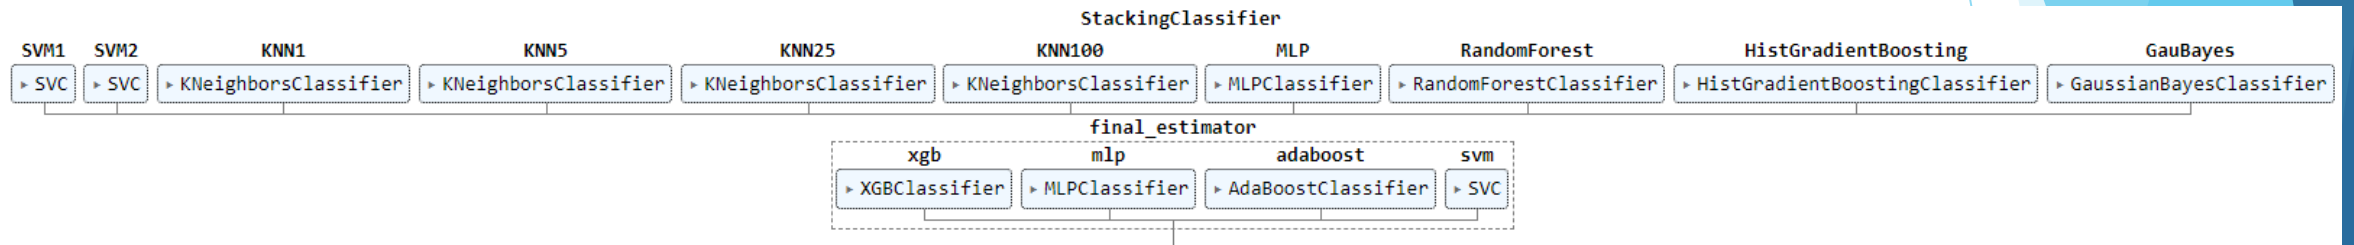
Model	Outer Fold Accuracy
SVM	84.14%
MLP	80%
1d-CNN	83.28%
XGBoost	75.44%
RandomForest	79.13%
HistGradBoosting	77.58%
Adaboost	71%
Stacking (meta = xgb)	84.12%
Stacking (meta = MLP)	86.01%
Stacking (meta = AdaBoost)	85.13%
Stacking (meta = SVM)	86.11%
Stacking (meta = {xgb, MLP, AdaBoost, SVM})	86.24%

# Part D: Observations

- ▶ General:
  - ▶ PCA, t-SNE and scaling did not improve results
- ▶ Stacking
  - ▶ Trying different parameters sets for the base models did not lead to improvement (due to limited computational resources not many sets were tried)
  - ▶ Fine tuning the meta-model using grid search provided improved performance in the outer loop
  - ▶ Fine tuning the models of the meta voting classifier did not lead to improved performance



# Part D: Final Model



- ▶ Model Tuning Observations:
- ▶ Fine-tuning base models increased time complexity significantly without noticeable improvement.
- ▶ Experimented with different meta-classifiers; soft voting classifier yielded the best results.
- ▶ No significant gains from fine-tuning the models within the voting classifier.

THE END